

## CS5530 Written Assignment #5

**Due: April 18, Wed, 11:30am.**

**Problem 1.** [40 pts]

• Database

- Employee-Table(ssn, name, salary)
  - \* E1(132, Smith, 20K)
  - \* E2(456, Kelley, 40K)
  - \* E3(678, Johnson, 400K)
  - \* E4(792, Preston, 40K)
  - \* E5(865, Johnson, 60K) ...
- DPT-table(dnumber, dname, budget)
  - \* D1(1, Marketing, 1M)
  - \* D2(2, Engineering, 2M)
  - \* D3(3, R&D , 4M)
  - \* D4(4, HR , 1M) ...

**Part 1.** Consider the following schedules

- 1) T1.R(E1), T1.W(E1), T2.R(E2), T2.W(E2), T1.R(D1), T1.commit, T2.commit
- 2) T1.R(E1), T1.W(E1), T2.R(E2), T2.W(E2), T2.W(E1), T1.R(E1), T1.commit, T2.commit
- 3) T1.R(E where salary>40K and salary<100k), T2.Insert(into E, (999, Bob, 50k)), T2.commit, T1.R(E where salary>40k and salary<100k), T1.commit
- 4) T1.R(E where salary>40K and salary<100k), T2.Insert(into E, (999, Bob, 50k)), T1.R(E where salary>40k and salary<100k), T2.commit, T1.commit.

For each schedule, please explain if it is valid in 1) the serializable isolation level; 2) the repeatable read isolation level.

**Part 2.** Consider the following scenario:

$T_1$	$T_2$
Begin Transaction  T1.Read(E1) T1.Update(set E1.salary=E1.salary*1.10) T1.Read(Select * from Employee)  T1.insert(Into E, (999, Bob, 50k)) T1.Read(Select * from Employee) T1.commit	Begin Transaction   T2.Update(set D4.budget=2M) T2.Delete(E5)
System Crash	

Figure 1: Sequence of events

Show the content of the table Employee after the system has recovered from the system crash.

**Problem 2.** [60 pts]

1. In the following schedules,  $R_i(A)$  stands for a Read(A) operation by transaction  $i$  and  $W_i(A)$  stands for a Write(A) operation by transaction  $i$ . For each of the following schedules show if it is conflict-serializable and give a conflict-equivalent serial schedule if it is one. Hint: use its dependency graph.

- (a)  $R_1(A), R_2(B), W_3(A), R_2(A), R_1(B), T1.Commit, T2.Commit, T3.Commit$   
 (b)  $R_1(A), R_2(B), W_1(A), R_3(C), W_2(B), W_3(C), R_4(D),$   
 $R_4(A), W_4(D), T1.Commit, T2.Commit, T3.Commit, T4.Commit$   
 (c)  $R_3(E), R_1(D), W_2(C), W_3(A), R_1(E), W_4(B), R_1(B), W_3(E),$   
 $R_4(A), W_4(C), T1.Commit, T2.Commit, T3.Commit, T4.Commit$

2. For the following 2 schedules, show if each is allowed in strict 2PL, and if not, what happens.

An example is given: (the schedule will not be allowed in strict 2PL as it is).

$S_1(A), R_1(A), X_2(A), W_2(A), X_1(B), R_1(B), W_1(B), T1.Commit, X_2(B), W_2(B), T2.Commit$

$T_1$	$T_2$		$T_1$	$T_2$
S(A)			S(A)	
R(A)			R(A)	
	X(A), Blocked			X(A), Blocked
X(B)			X(B)	
R(B)			R(B)	
W(B)			W(B)	
Commit		OR	Commit	
Release S(A)			Release S(A)	
	W(A)		Release X(B)	
	X(B), Blocked			W(A)
Release X(B)				X(B)
	W(B)			W(B)
	Commit+Release its locks			Commit+Release its locks

- (a)  $S_1(A), R_1(A), S_2(B), R_2(B), S_3(C), R_3(C), X_3(D), W_3(D), T3.Commit,$   
 $X_2(C), W_2(C), T2.Commit, X_1(B), W_1(B), T1.Commit$   
 (b)  $X_1(A), R_1(A), X_2(B), R_2(B), X_3(C), R_3(C), S_1(B), R_1(B), S_2(C), R_2(C),$   
 $S_3(A), R_3(A), W_1(A), T1.Commit, W_2(B), T2.Commit, W_3(C), T3.Commit$

3. For the following schedules, show what happens in each case.

An example of non-strict 2PL is given (either is fine; in fact, there could be even more. But showing one is enough):

$S_1(A), R_1(A), X_1(B), X_2(A), W_2(A), R_1(B), W_1(B), T1.Commit, X_2(B), W_2(B), T2.Commit$

$T_1$	$T_2$		$T_1$	$T_2$
S(A)			S(A)	
R(A)			R(A)	
X(B)			X(B)	
	X(A), Blocked		Release S(A)	
Release S(A)				X(A)
	W(A)			W(A)
R(B)		OR	R(B)	
W(B)			W(B)	
Release X(B)			Release X(B)	
Commit			Commit	
	X(B)			X(B)
	W(B)			Release X(A)
	Release X(A), Release X(B)			W(B)
	Commit			Release X(B)
				Commit

- (a)  $X_1(B), W_1(B), X_2(A), W_2(A), S_2(B), R_2(B), S_1(A), R_1(A), T1.Commit, T2.Commit$ .

Using **strict 2PL**

- (b)  $X_1(B), W_1(B), X_2(A), W_2(A), S_2(B), R_2(B), S_1(A), R_1(A), T1.Commit, T2.Commit$ .

Using **non-strict 2PL**

- (c)  $X_1(B), W_1(B), S_1(A), S_2(B), R_2(B), R_1(A), X_2(A), W_2(A), T1.Commit, T2.Commit$ .

Using **non-strict 2PL**